

1 General Definitions

Definitions:

- **Cloud Operations** is the practice of managing and optimizing cloud-based services and infrastructure.
- **GitOps:** Git-based infrastructure and application deployment; uses Git as single source of truth; enables CI/CD, automation, version control, and declarative configuration.
- **DevOps** combines development and operations; focuses on automation, collaboration, CI/CD, monitoring, and agile delivery.

1.1 DevOps Cycle

Plan (add Objectives and Requirements to Backlog), **Code** (add Code to Repo), **Build** (Pipelines runs on push, builds and unit tests software), **Test** (Build is deployed to staging environment, tested using E2E, load, accessibility tests), **Release** (snapshot of code is versioned, changes are documented), **Deploy** (release is installed onto production environment), **Operate** (application should run smoothly, issues are troubleshooted and documented, infrastructure is scaled), **Monitor** (Application Data is gathered and used for planning) **Difference Between Continuous Delivery & Continuous Deployment:** Deployment automatically pushes from staging to production, in Delivery this is manual. **CD&D Deployment Strategies:**

- *Rolling Deployment:* Update infrastructure gradually, minimal downtime
- *Blue-Green:* Two environments: Old and new versions respectively
- *Canary:* Small user group tests first
- *Feature Flag:* Deploy but activate later, can be toggled
- *Dark Launching:* Rolling out a feature invisible for users, test its performance in the background

2 GitLab

Example GitLab pipeline:

```
stages:
  - build
  - test
  - deploy
cache:
  paths:
    - .cache/
build:
  stage: build
  script:
    - echo "Building..."
    - mkdir -p artifacts && echo "artifact" > artifacts/output.txt
artifacts:
  paths:
    - artifacts/
  expire_in: 1 hour
```

```
test:
  stage: test
  dependencies:
    - build
  script:
    - cat artifacts/output.txt
deploy_staging:
  stage: deploy
  environment:
    name: staging
    url: https://staging.example.com
    on_stop: stop_staging # Unstages the env
  script:
    - cat k8.yaml | envsubst | kubectl
    apply -f -
artifacts:
  expire_in: 1 hour
stop_staging:
  stage: deploy
  environment:
    name: staging
  action: stop
  script:
    - echo "Stopping staging"
```

2.1 Environments

Describe where the code gets deployed (e.g. Local, Integration, Testing, Staging, Production). Can be linked to a K8 cluster (needs to be set up via GitLab UI):

2.2 Push- vs. Pull-Based Deployments

Push-Based: + Easy to use, + flexible deployment targets, - firewall needs to be opened, - pipeline needs to be adjusted for new environments **Pull-Based:** + no need for open firewall, + better scaling, - agent needs to be installed in every cluster

3 Terraform

TF doesn't speak directly with an SDK, but rather Terraform -> Provider -> Client SDK. Different providers enable different platforms (AWS, Azure, Kubernetes, ...). A sample in HCL (Hashicorp Configuration Language):

```
variable "instance_type" {
  default = "t2.micro"
}
provider "aws" {
  region = "us-east-1"
}
resource "aws_instance" "web" {
  ami = "ami-0c55b159c9bfafe1f0"
  instance_type = var.instance_type
}
output "public_ip" {
  value = aws_instance.web.public_ip
}
```

To deploy infrastructure, write HCL in files like main.tf, then run terraform init, terraform plan to show changes that would be made, terraform apply to actually apply the changes. Use terraform destroy to delete all made changes.

3.1 State

Terraform stores state in the terraform.tfstate file. When working in teams, this state file also has to be shared as the terraform command

relies on is validity. This could for example be done via an S3 Bucket.

4 Ansible

Ansible can be used to provision servers. It does not have statefiles and is idempotent, meaning it won't make changes unless it has to.

4.1 Infrastructure

In a network of servers, one server is the **host**. The host can connect to other machines using SSH. On the host, playbooks can be written in yaml files. Run a playbook by using **ansible-playbook playbook.yaml**

```
- name: Example Playbook
hosts: web
become: true
vars:
  packages:
    - nginx
    - curl
  enable_service: true
  secret_password: "{{ vault_password }}"
roles:
  - myrole
tasks:
  - name: Install packages
    apt:
      name: "{{ item }}"
      state: present
    loop: "{{ packages }}"
    notify: restart nginx
  - name: Configure app if enabled
    template:
      src: app.conf.j2
      dest: /etc/app.conf
    when: enable_service
    tags: config
handlers:
  - name: restart nginx
    service:
      name: nginx
      state: restarted
```

4.2 Vaults

Vaults can be used to encrypt data: The file vault.yaml with the contents vault_password: \$supersecret" can be encrypted using ansible-vault encrypt vault.yml and then included in a play: ansible-playbook playbook.yml --ask-vault-pass To create a file, use ansible-vault create foo.yaml

4.3 Collections, Roles & Tags

Collections are bundles of plugins, roles and modules. Install them using ansible-galaxy collection install <name>, or define a requirements.yaml to install multiple collections at once. **Roles** are an abstraction above playbooks, allowing to reuse configuration steps: create a role using ansible-galaxy init <name>, then use a role like in the example above. **Tags** can be used to execute a subset of tasks instead of the whole playbook. Run only specific tags by appending -tags <name> at the end of the ansible-playbook command. There are also two special commands: Tag

always runs every time, except when explicitly skipped: -skip-tags=always. Tag never does not run unless specified with -tags=never

4.4 Jinja2

Jinja2 is the templating engine which is used by Ansible. It is used to generate configuration files.

5 Kubernetes (K8)

K8 Objects: Persistent entities which signal an intent (e.g. for something to be created on the cluster) **K8 Controller:** Tracks a Object and is responsible for bringing the current State closer to the desired State.

Pod: Represents a process running on your cluster. Should contain one container, multiple are possible.

Sidecars: Sidecars are containers that run along the primary container in the same pod. Example use cases might be logging, security, data synchronization.

Init Containers: Similar to sidecars, but run and finish before app containers.

Volume: Assigns physical Storage to a Pod

ReplicaSet: Makes sure that a specified number of replica pods are running. In practice, deployments are used.

Deployment: Allows to manage one or multiple Pods.

Service: API Resource to expose logical set of Pods in the namespace. Acts as a load balancer (round-robin).

Ingress: Provides external Access to a Service.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app
spec:
  replicas: 3 # automatically deploys replicaset
  selector:
    matchLabels:
      app: web
  template:
    metadata: # pod
      labels:
        app: web
    spec:
      strategy:
        type: RollingUpdate # rolling update
      rollingUpdate:
        maxUnavailable: 1
        maxSurge: 1
      initContainers:
        - name: init
          image: busybox
          command: ['sh', '-c', 'sleep 10']
      containers: # container
        - name: web
          image: nginx
          ports:
            - containerPort: 80
        - name: sidecar
          image: busybox
          command: ['sh', '-c', 'while true; do sleep 30; done']
```

5.1 Namespaces

Used to separate resources. Only resources in same namespace can communicate directly.

5.2 Rolling Updates

Rolling updates can be used in order to ensure that enough pods are always running. Rolling updates can be using maxUnavailable (maximum No. of Pods upgrading at the same time) and maxSurge (max No. of Pods allowed to run beyond specified No. of replica)

```
spec:
```

5.3 Scheduling

The kube-scheduler determines which nodes run which Pods. We can influence this decision process:

```
kind: Pod
spec:
  nodeSelector:
    disktype: ssd # this label needs to be in pod.spec
```

To evaluate if a Node is eligible to run a Pod, the following things are considered: Port availability, CPU & Memory resources, available volumes, specified labels. Additionally, scoring is used to evaluate remaining nodes with criteria: pods of same service should be on different nodes, nodes with few used resources are prioritized, node affinity.

Taints can also be applied on nodes and pods, pods won't be deployed on nodes with matching taints. Tolerations can be used to make exceptions to taints. Types of taints: NoSchedule, PreferNoSchedule, NoExecute

5.4 Commands

Apply a manifest.yaml: kubectl create|apply|replace -f manifest.yaml

Connecting to a Pod: kubectl exec -it nginx-xxx - sh **Undo rollout:** kubectl rollout undo

6 Helm

A package manager for K8, enabling to reuse configurations for common use cases (DB, monitoring). Helm provides *Charts*, which are a collection of yaml files describing different K8 Objects. When deploying a chart on your cluster, it is called a *Release*. Charts are available through different *Repositories*.

Helm charts use templating (like {{Release.Name}}), for information about package, or {{.Values.xyz.abc | default example}} for information passed by values.yaml or via commandline -set)

6.1 Commands

```
# Repo commands
```

```
helm repo list
helm repo add <repo> <url>
helm repo rm <repo>
```

```
helm list # list installed releases
```

```
helm install -f values.yaml <release> <chart> # install with custom values
```

```
helm show <chart> | readme | values | all>
<chartname> | repo>
```

```
helm upgrade <release> <chart>
helm rollback <release> <revision>
```